


Behavioral sufficiency: does the model SAY it?



Causal influence: does it PREFER it?

